

Synthesis

TP0

Preliminary

- Relational Database Management Systems (RDBMS)
 - Supports transactions with ACID properties
 - Atomicity: A transaction block is either fully executed or completely canceled
 - Consistency (or Correctness): The resulting database is valid w.r.t to the integrity constraints.
 - Isolation: The effect of two concurrent transactions is the same as if one was scheduled before the other.
 - Durability: Once confirmed, a transaction cannot be rolled back.
 - Inefficient for sparse relations
 - Not suitable for aggregations
 - Need query optimization
 - Queries can be time expensive
- Relational model
 - Tables or relations, each of which have several columns or attributes, each of which have a type
 - Data is stored as records or tuples
 - Sets
 - \mathcal{L} of labels
 - \mathcal{V} of values
 - \mathcal{T} of types (for each $\tau \in \mathcal{T}$, $\tau \subseteq \mathcal{T}$)
 - Relation schema is a n-tuple (A_1, \dots, A_n) where each A_i is a pair (L_i, τ_i) with $L_i \in \mathcal{V}$ and $\tau_i \in \mathcal{T}$

Questions

- In this zeroth exercise, we will query a single the following tables (with its fields):
 - Room (Name, Time, Movie Title)
 - Movie (MovieTitle, Director, Actor)
 - Producer (ProducerName, MovieTitle)
 - Seen (Spectator, MovieTitle)

- Like (Spectator, MovieTitle)
- 1: Where and when can we see the movie "Mad Max...?"
`DROP(FILTER(Room, movieTitle="Mad Max"), MovieTitle)`
- 2: What are the movies directed by Welles?
`DROP(DROP(FILTER(Movie, Director="Welles"), Director), Actor)`
- 3: Who are the actors of "Ran"?
`DROP(DROP(FILTER(Movie, MovieTitle="Ran"), Director), MovieTitle)`
- 4: Where can we see a movie in which Signoret plays?
`moviesWithS = DROP(DROP(FILTER(Movie, Actor="Signoret"), Actor), Director)`
`whereMovie = RENAME(DROP(Movie, Time), MovieTitle, MovieTitleBis)`
`DROP(FILTER(PRODUCT(moviesWithS, whereMovie), movieTitle=movieTitleBis), [MovieTitle, MovieTitleBis])`
- 5: Among the actors who produced at least one movie?
`actors = DROP(DROP(Movie, MovieTitle), Director)`
`producers = DROP(Producers, MovieTitle)`
`DROP(FILTER(PRODUCT(actors, producers), ProducerName=Actor), ProducerName)`
- 6: Among the actors who directed a movie that they played in?
`DROP(FILTER(Movie, Director=Actor), [MovieTitle, Director])`
- 7: Who plays in one (or more) movie from Varda?
`DROP(FILTER(Movie, Director="Varda"), [MovieTitle, Director])`
- 8: Who are the actors playing in all the films from Chloé Zhao?
`ZhaoMovies = DROP(FILTER(Movie, Director="Zhao"), [Director, Actor])`
`allPairsA_ZM = PRODUCT(actors, ZhaoMovies)`
`missingPairs = DIFFERENCE(allPairsA_ZM, Movie)`
`actorsInAllZhao = DIFFERENCE(actors, missingPairs)`

- 9: Who produces all the movies from Kurosawa?
`KurosawaMovies = DROP(FILTER(Movie, Director="Kurosawa"), [Director, Actor])`
`AllPairsKM_P = PRODUCT(producers, KurosawaMovies)`
`missingPairs = DIFFERENCE(Producer, AllPairsKM_P)`
`producersNotInAllK = DIFFERENCE(Producer, missingPairs)`
`DIFFERENCE(Producer, producersNotInAllK)`
- 10: Who are the spectators watching all the movies?
`movieNames = DROP(Movie, [Director, Actor])`
`spectators = DROP(Seen, MovieTitle)`
`allPairs_ms = PRODUCT(movieNames, spectators)`
`missingPairs = DIFFERENCE(allPairs_ms, Seen)`
`DIFFERENCE(spectators, missingPairs)`
- 11: Among the spectators, who like all the movies they see?
`DIFFERENCE(Seen, DIFFERENCE(Seen, Like))`
- 12: Where can we see Adèle Haenel after 16:00?
`moviesRoomAfter16 = DROP(FILTER(Room, time>16:00), Time)`
`moviesWithHaenel = RENAME(DROP(FILTER(Movie, Actor=Haenel), [Director, Actor]), MovieTitle, MovieTitleBis)`
`DROP(FILTER(PRODUCT(moviesRoomAfter16, moviesWithHaenel), movieTitle=movieTitleBis), [movieTitle, movieTitleBis])`
- 13: What are the movies with no room projecting them?
`DIFFERENCE(Movie, Rooms)`
- 14: Among the producers, who produce movies shown nowhere?
`DIFFERENCE(Producer, Room)`
- 15: Among the producers who saw all the movies they directed?
`AllPairs_P_M = PRODUCT(movieNames, producers)`
`missingPairs = DIFFERENCE(AllPairs_P_M, Producer)`
`DIFFERENCE(producers, missingPairs)`

- 16: Among the spectators, who saw all the movies from Kurosawa
`moviesK = DROP(FILTER(Movie, Director="Kurosawa"), [Director, Actor])`
`allPairs_ms = PRODUCT(moviesK, spectators)`
`missingPairs = DIFFERENCE(allPairs_ms, Seen)`
`DIFFERENCE(spectators, missingPairs)`
 - 17: Among the spectators, who liked a movie they did not watch?
`DROP(DIFFERENCE(Like, Seen), MovieTitle)`
 - 18: Among the spectators, who liked 0 movies?
`DIFFERENCE(DROP(Seen, MovieTitle), DROP(Like, MovieTitle))`
- ## TP1
- ## Preliminary
- SQL types
 - BOOLEAN
 - INT: Signed range is from -2147483648 to 2147483647
 - SERIAL for an auto-incrementing identifier
 - REAL for floating-point numbers (4-byte)
 - NUMERIC for high-precision numbers (1000 digits)
 - TEXT or VARCHAR: text
 - BLOB for binary strings: up to 65,535 bytes
 - TIMESTAMP for date and time
 - The subquery must have an alias (here, M1), even if it is not used
- ## Questions
- In this first exercise, we will query a single world table. The table contains information about the world's countries. The table has the following fields:
 - name varchar(50): the English name of the country (primary key)
 - continent
 - varchar(60): the English name of the continent
 - area decimal(10,0): the area in square kilometers (note: it is sometimes missing)
 - population decimal(11,0): the population (note: it is sometimes missing)
 - capital varchar(60): the English name of the capital of the country

Basics: countries and continents

- 1: Modify the query below to select the name of all countries in Europe. Hint: use the WHERE clause:

```
SELECT
  name
FROM
  world
WHERE
  continent = 'Europe';
```
- 2: Select all DISTINCT continents in the database:

```
SELECT
  DISTINCT continent
FROM
  world;
```

Filtering names

- 3: Select all country names that start with an 'F'. Hint: use LIKE. Note: with PostgreSQL, the LIKE operator is case-sensitive.

```
SELECT
  name
FROM
  world
WHERE
  name LIKE 'F%';
```
- 4: Select all country names containing the letter 'z':

```
SELECT
  name
FROM
  world
WHERE
  name LIKE '%z%';
```

Population

- 5: Select all country names having at most one million inhabitants:

```
SELECT
  name
FROM
  world
WHERE
  population <= 1000000;
```
- 6: Write a query returning the name and population of all European countries

```
SELECT
  name,
  population
FROM
  world
WHERE
  continent = 'Europe';
```
- 7: Write a query returning the country name and population of France, Germany, and Italy (in that order):

```
SELECT
  name,
  population
FROM
  world
WHERE
  name IN ('France', 'Germany', 'Italy')
```

ORDER BY name;

- 8: Write a query returning the population of France:

```
SELECT
  population
FROM
  world
WHERE
  name = 'France';
```

Boolean conditions

- 9: Write a query selecting the name, population, and area of all countries that have at least 100 million inhabitants or have an area of at least 3 million square kilometers (or both). Order the results by name.

```
SELECT
  name,
  population,
  area
FROM
  world
WHERE
  population >= 100000000
OR area >= 3000000
ORDER BY
  name;
```

- 10: Write a query selecting the name, population, and area of all European countries which have population greater than 50 million or area greater than 500000 square kilometers (or both), ordered by name. Hint: be careful!

```
SELECT
  name,
  population,
  area
FROM
  world
WHERE
  continent = 'Europe'
AND (
  population >= 50000000
OR area >= 500000
)
ORDER BY
  name;
```

- 11: Write a query selecting the name, population, and area of all countries that are either large in population or large in size, but not both. Specifically, select countries that have at least 100 million inhabitants OR have an area of at least 3 million square kilometers, but not both. Order the results by name

```
SELECT
  name,
  population,
  area
FROM
  world
WHERE
  (
    population >= 100000000
    AND NOT area >= 3000000
  )
OR (
    area >= 3000000
    AND NOT population >=
```

```

100000000
)
ORDER BY
name;

```

String operations

- 12: Write a query selecting the country names that are the same as that country's capital, sorted by name.

```

SELECT
name
FROM
world
WHERE
capital = name
ORDER BY
name;

```

- 13: Write a query selecting the country names and capital names where both names have the same length. E.g., "Greece" and "Athens" both have length 6. Exclude the cases from the previous question, i.e., those where the country name is the same as the capital, and sort the results by country name. Hint: search the Internet for a function to compute the length of a string

```

SELECT
name,
capital
FROM
world
WHERE
CHAR_LENGTH(capital) =
CHAR_LENGTH(name)
AND capital <> name
ORDER BY
name;

```

- 14: Write a query returning the country name with the least population, along with its population.

```

SELECT
name,
population
FROM
world
ORDER BY
population
LIMIT
1;

```

- 15: Write a query returning the name and population of the 5 countries with the greatest population (ordered by descending population, i.e., from greatest to least population), where this population is known. Note: use IS NOT NULL.

```

SELECT
name,
population
FROM
world
WHERE
population IS NOT NULL
ORDER BY
population DESC
LIMIT
5;

```

- 16: Write a query returning the name, population, and area of the 100th country in alphabetical order

```

SELECT
name,
population,
area
FROM
world
ORDER BY
name
LIMIT
1 OFFSET 99;

```

- 17: Write a query returning the name and population of the 10 most populous countries among the 20 countries with the greatest area. Hint 1: remember that the population and the area may be NULL! Hint 2: use a subquery, and remember to give an alias to its result!

```

SELECT
*
FROM
(
SELECT
name,
population
FROM
world
WHERE
area IS NOT NULL
AND population IS NOT
NULL
ORDER BY
area DESC
LIMIT
20
) AS T
ORDER BY
population DESC
LIMIT
10;

```

Computation and aggregation

- 18: Write a query returning the name and population density of the Asian countries. The population density is the population divided by the area. Order the results by name.

```

SELECT
name,
population / NULLIF(area,
0) AS density
FROM
world
WHERE
continent = 'Asia'
ORDER BY
name;

```

- 19: Write a query returning the name of Asian countries with a new column dense containing "yes" for countries with density at least 100 inhabitants per square kilometer, and "no" otherwise. Sort the results to have first the dense countries, then the non-dense countries, and then order the countries in each group by name. Hint: use UNION.

```

SELECT
name,
'yes' AS dense
FROM

```

```

world
WHERE
continent = 'Asia'
AND population /
NULLIF(area, 0) >= 100
UNION
SELECT
name,
'no' AS t
FROM
world
WHERE
continent = 'Asia'
AND population /
NULLIF(area, 0) < 100
ORDER BY
dense DESC,
name;

```

```

SELECT
name,
CASE
WHEN (density > 100)
THEN 'yes'
ELSE 'no'
END AS dense
FROM
(
SELECT
name,
population / area AS
density
FROM
world
WHERE
(
area IS NOT NULL
AND area != 0
AND continent =
'Asia'
)
ORDER BY
name
);

```

- 20: Write a query returning the number of countries in the database. Hint: use COUNT(*).

```

SELECT
COUNT(*)
FROM
world;

```

- 21: Write a query returning the total world population. Hint: use SUM.

```

SELECT
SUM(population)
FROM
world;

```

- 22: Write a query returning the continent names and, for each continent, the number of countries in that continent. Hint: use aggregation. In this question and the next ones, unless otherwise stated, order the results by continent.

```

SELECT
continent,
COUNT(*)
FROM
world
GROUP BY
continent
ORDER BY
continent;

```

- 23: Write a query returning the continent names and, for each continent, the total population and total area of the countries in that continent.

```

SELECT
continent,
SUM(population),
SUM(area)
FROM
world
GROUP BY
continent
ORDER BY
continent;

```

- 24: Show the name and total population of continents having total population at least 100 million.

```

SELECT
continent,
SUM(population) AS total
FROM
world
GROUP BY
continent
HAVING
SUM(population) >=
100000000
ORDER BY
continent;

```

- 25: Show the name of continents and the number of countries in that continent that have population at least 1 million.

```

SELECT
continent,
COUNT(*)
FROM
world
WHERE
population >= 1000000
GROUP BY
continent
ORDER BY
continent;

```

- 26: Compute the average population of the world's countries.

```

SELECT
AVG(population) AS average
FROM
world;

```

- 27: Compute, for every continent name, three columns containing the following three values, rounded to the nearest integer with the ROUND function. Do you understand why the three values are different?
 - its population density (i.e., the continent's population divided by the continent's area, taking all countries into account),
 - the average of the population densities of its countries (for the countries with a non-NULL and non-zero area),
 - the sum of the population densities of its countries (for the countries with a non-NULL and non-zero

```

area), divided by the
number of countries
SELECT
continent,
ROUND(SUM(population) /
SUM(area)) AS total,
ROUND(AVG(population /
NULLIF(area, 0))) AS
average,
ROUND(SUM(population /
NULLIF(area, 0)) / COUNT(*))
AS fake_average
FROM
world
GROUP BY
continent;

```

- 28: Return a table with a column alpha containing the first letter of country names (ordered alphabetically) and a column total with the total population of countries whose name starts with that letter. Hint: use SUBSTR as in the example.

```

SELECT
SUBSTR(name, 1, 1) AS
alpha,
SUM(population)
FROM
world
GROUP BY
alpha
ORDER BY
alpha;

```

- 29: Compute, for every continent name, its country having the largest area, and its country having the greatest population. Order the result by continent. Hint: this is a complicated task, use subqueries and/or join the world table with itself. Note: It is normal that Kazakhstan appears as a country in Europe.

```

SELECT
W1.continent,
W1.name AS largest,
W2.name AS most_populous
FROM
world AS W1,
world AS W2
WHERE
W1.continent =
W2.continent
AND W1.area IS NOT NULL
AND W2.population IS NOT
NULL
AND NOT EXISTS (
SELECT
1
FROM
world AS W
WHERE
W.continent =
W1.continent
AND W.area > W1.area
)
AND NOT EXISTS (
SELECT
1
FROM
world AS W
WHERE
W.continent =
W2.continent
AND W.population >
W2.population

```

```

)
ORDER BY
W1.continent;

```

- 30: Compute, for every continent name, a count of how many countries in the continent are strictly more populous than the continent's largest country (by area). Order the results by continent, and do not omit results where the count is zero.

```

SELECT
continent,
SUM(total)
FROM
(
SELECT
continent,
COUNT(*) AS total
FROM
world AS W1
WHERE
population > (
SELECT
population
FROM
world AS W2
WHERE
W1.continent =
W2.continent
AND W2.area IS NOT
NULL
ORDER BY
area DESC
LIMIT
1
)
) GROUP BY
continent
UNION
SELECT
DISTINCT continent,
0 AS TOTAL
FROM
world
) AS T
GROUP BY
continent
ORDER BY
continent;

```

TP2

Questions

- In this second exercise, we will query three tables: movie, casting, and actor.
 - The movie table describes movies, and contains the following fields:
 - id int(11), an identifier
 - title varchar(50), the title of the movie
 - yr int(11), the year the movie was released
 - director int(11), the identifier of the director
 - The actor table contains people (actors and directors) and contains the following fields:
 - id int(11), an identifier

- name varchar(50), the name of the person
- The casting describes actors starring in movies. It contains the following fields:
 - movieid int(11), the identifier of the movie in the movie table
 - actorid int(11), the identifier of the actor in the actor table
 - ord int(11), an integer describing the position of the actor in the film's starring list. The first actor (called the leading actor) has ord value 1, the second actor has ord value 2, and so on.

Basic joins

- 1: Compute the title of every Star Wars movie (starting with "Star Wars") and the name of its director. Sort the result by title.

```
SELECT
  title,
  name
FROM
  movie,
  actor
WHERE
  title LIKE 'Star Wars%'
  AND director = actor.id
ORDER BY
  title;
```

- 2: Compute the list of the names of the actors starring in the movie "Jurassic Park" (1993), in the order in which they starred (i.e., by increasing ord value).

```
SELECT
  name
FROM
  movie,
  actor,
  casting
WHERE
  movie.id = casting.movieid
  AND actor.id =
  casting.actorid
  AND title = 'Jurassic
  Park'
ORDER BY
  ord;
```

- 3: Compute the list of the titles of the movies where "George Clooney" appeared, ordered by title.

```
SELECT
  title
FROM
  movie,
  actor,
  casting
WHERE
  movie.id = casting.movieid
  AND actor.id =
  casting.actorid
  AND name = 'George
  Clooney'
ORDER BY
  title;
```

- 4: Compute the list of the titles of all movies released in 1920 together

with the name of their leading actor (the one with ord value of 1), ordered by title.

```
SELECT
  title,
  name
FROM
  movie,
  casting,
  actor
WHERE
  yr = 1920
  AND casting.movieid =
  movie.id
  AND casting.actorid =
  actor.id
  AND ord = 1
ORDER BY
  title;
```

- 5: Compute the list of the titles of all movies released in 1920 together with the name of their leading actor and with the name of their director, ordered by title.

```
SELECT
  title,
  SA.name,
  D.name
FROM
  movie,
  casting,
  actor AS SA,
  actor AS D
WHERE
  yr = 1920
  AND casting.movieid =
  movie.id
  AND casting.actorid =
  SA.id
  AND ord = 1
  AND D.id = director
ORDER BY
  title;
```

- 6: Compute the five movies in the database having the highest number of participating actors, and this number of participating actors, sorted by decreasing number of actors. Warning: beware of titles like "The Hunchback of Notre Dame" that are the titles of multiple movies!

```
SELECT
  title,
  COUNT(*) AS cnt
FROM
  movie,
  casting
WHERE
  casting.movieid = movie.id
GROUP BY
  movie.id,
  title
ORDER BY
  cnt DESC
LIMIT
  5;
```

- 7: Compute the years where the actor "Rock Hudson" participated to strictly more than one movie, along with the number of movies to which he participated on that year, sorted by decreasing number of movies, then by ascending year.

```
SELECT
  yr,
  COUNT(*) AS count
FROM
  movie,
  casting,
  actor
WHERE
  movie.id = casting.movieid
  AND actor.id =
  casting.actorid
  AND name = 'Rock Hudson'
GROUP BY
  yr
HAVING
  COUNT(*) > 1
ORDER BY
  COUNT(*) DESC,
  yr;
```

- 8: Compute the names of actors who were the leading actor in a movie where Harrison Ford appeared (and were not Harrison Ford himself). Order the results by actor name.

```
SELECT
  DISTINCT A2.name
FROM
  actor AS A1,
  actor AS A2,
  casting AS C1,
  casting AS C2
WHERE
  A1.name = 'Harrison Ford'
  AND A1.id = C1.actorid
  AND C1.movieid =
  C2.movieid
  AND C2.actorid = A2.id
  AND C2.ord = 1
  AND A2.name <> 'Harrison
  Ford'
ORDER BY
  A2.name;
```

- 9: Compute the titles of movies which were both directed by Woody Allen and had Woody Allen appear as an actor, sorted in alphabetical order.

```
SELECT
  title
FROM
  movie,
  actor,
  casting
WHERE
  casting.movieid = movie.id
  AND casting.actorid =
  actor.id
  AND actor.name = 'Woody
  Allen'
  AND movie.director =
  actor.id
ORDER BY
  title;
```

- 10: Compute the titles of movies which were directed by Woody Allen or had Woody Allen appear as an actor (or both), sorted in alphabetical order.

```
SELECT
  title
FROM
  movie,
  actor,
  casting
WHERE
```

```
casting.movieid = movie.id
  AND casting.actorid =
  actor.id
  AND actor.name = 'Woody
  Allen'
UNION
SELECT
  title
FROM
  movie,
  actor
WHERE
  movie.director = actor.id
  AND actor.name = 'Woody
  Allen'
ORDER BY
  title;
```

Trick questions

- 11: In which movie title did Alain Delon and Catherine Deneuve appear together?

```
SELECT
  title
FROM
  movie,
  casting AS C1,
  casting AS C2,
  actor AS A1,
  actor AS A2
WHERE
  A1.id = C1.actorid
  AND A2.id = C2.actorid
  AND C1.movieid = movie.id
  AND C2.movieid = movie.id
  AND A1.name = 'Alain
  Delon'
  AND A2.name = 'Catherine
  Deneuve';
```

- 12: Find the only actor who appeared in all Star Wars movies.

```
SELECT
  name
FROM
  actor
WHERE
  NOT EXISTS (
    SELECT
      id
    FROM
      movie
    WHERE
      TITLE LIKE 'Star
      Wars%'
      AND NOT EXISTS (
        SELECT
          +
        FROM
          casting
        WHERE
          movieid = movie.id
          AND actorid =
          actor.id
      )
  );
```

- 13: Find the only actor which only appeared in movies where Harrison Ford appeared, and appeared in strictly more than one such movie.

```
SELECT
  name
FROM
  actor,
  casting,
  movie
WHERE
```

```
casting.actorid = actor.id
  AND casting.movieid =
  movie.id
  AND name <> 'Harrison
  Ford'
AND NOT EXISTS (
  SELECT
    1
  FROM
    movie AS M1,
    casting AS C1
  WHERE
    C1.movieid = M1.id
    AND C1.actorid =
    actor.id
    AND NOT EXISTS (
      SELECT
        1
      FROM
        casting AS C2,
        actor AS A2
      WHERE
        A2.name =
        'Harrison Ford'
        AND C2.actorid =
        A2.id
        AND C2.movieid =
        C1.movieid
      )
  )
GROUP BY
  name
HAVING
  COUNT(movie.id) > 1;
```

- 14: For performance reasons, we limit ourselves to the movies released no later than 1930, and to the actor names starting with A, B, or C. We say that two actors X and Y are challengers if X was the leading actor in a movie where Y appeared and Y was the leading actor in a movie where X appeared. Compute all pairs X, Y of challengers (with X < Y, in alphabetical order of X and then of Y).

```
WITH oldmov AS (
  SELECT
    id,
    yr
  FROM
    movie
  WHERE
    yr <= 1930
),
aactor AS (
  SELECT
    id,
    name
  FROM
    actor
  WHERE
    name LIKE 'A%'
    OR name LIKE 'B%'
    OR name LIKE 'C%'
)
SELECT
  DISTINCT A1.name,
  A2.name
FROM
  aactor AS A1,
  aactor AS A2,
  casting AS C1a,
  casting AS C1b,
  casting AS C2a,
  casting AS C2b,
  oldmov AS Ma,
  oldmov AS Mb
WHERE
  A1.id = C1a.actorid
```

```
AND A1.id = C1b.actorid
  AND A2.id = C2a.actorid
  AND A2.id = C2b.actorid
  AND C1a.movieid =
  C2a.movieid
  AND C1b.movieid =
  C2b.movieid
  AND C1a.ord = 1
  AND C2b.ord = 1
  AND C1a.movieid = Ma.id
  AND C1b.movieid = Mb.id
  AND A1.name < A2.name
ORDER BY
  A1.name,
  A2.name;
```

TP3

Preliminary

- A relation satisfies the First Normal Form (1NF) if the data of every cell is an atomic type
 - There are only Single Valued Attributes
 - Attribute Domain does not change
 - There is a unique name for every Attribute/Column
 - The order in which data is stored does not matter
- Functional dependency (FD) on a relation R is an assertion of the form $A_1 \dots A_n \rightarrow B_1 \dots B_m$ where A_i and B_j are attributes of R
 - Constraint that always holds
 - Always holds if
 - $\{B_1 \dots B_m\} \subseteq \{A_1 \dots A_n\}$ (trivial FD)
 - at least $A_1 \dots A_n$ are in the left-hand side if $A_1 \dots A_n$ are a key
 - $A_1 \dots A_n \rightarrow B_1 \dots B_m$ if $A_1 \dots A_n \rightarrow B_j$ is an FD for each B_j
- A relation is in Boyce-Codd Normal Form (BCNF) if for every non-trivial FD $A_1 \dots A_n \rightarrow B_1 \dots B_m$ that it satisfies, then $A_1 \dots A_n$ is a superkey
 - Disallows
 - FDs between non-key attributes (attributes outside the key)
 - FDs from a strict subset of the key attributes
 - Non-BCNF
 - Many-to-many relationship
- Schema design
 - Being complete, i.e., can represent everything that is needed
 - Being clear to developers and as simple as possible

- Being precise: clear how to map actual business needs to data
- Not being too broad, i.e., correctly reflect constraints that are assumed
- Avoiding redundancy: make sure every data item is in one place
- Ensuring good performance (often linked to simplicity)
- Entity-relationship diagrams
 - Attribute
 - Single-valued attributes
 - Multi-valued attributes
 - Derived attributes (dashed circle) are derivable from stored attributes
 - Key of weak entities (dash-underlined)
 - Entity
 - Weak entity (double lined square)
 - Relationship
 - Identifying relationships (double lined diamond)
 - is-A Δ
 - Specialization: top-down
 - Generalization: bottom-up
 - Can be entities for other relationships
 - Roles are written above the lines connecting the relationship and entity
 - Cardinality constraints
 - One-to-one: functional and injective
 - One-to-many: injective but not functional
 - Many-to-one: functional but not injective
 - Many-to-many arbitrary and not functional
 - Below the role the minimal and maximal number of relationships in which an entity can participate (* means no limit)

Questions

PostgreSQL

- List databases \l, \l+
 - \x on for narrower settings
- Connect to a database \c < databaseName >
- Display tables \dt
- Display schema \d < tableName >
- Display user roles \du
- Create a database \$ createdb -U postgres < databaseName >
- Create a database \$ dropdb -U postgres < databaseName >
- To alter a table ALTER TABLE < table > ...
 - RENAME to < newName >;
 - ADD < columnName > < dataType >;
 - DROP < columnName >;
 - RENAME < oldColumnName > to < newColumnName >;
 - ALTER < columnName > < dataType >;

Schema and views

- Schema


```
CREATE TABLE IF NOT EXISTS player (
  id SERIAL,
  name TEXT NOT NULL,
  money INT CHECK (money >= 0),
  PRIMARY KEY (id)
);

CREATE TABLE IF NOT EXISTS type (
  id SERIAL,
  name TEXT NOT NULL,
  PRIMARY KEY (id)
);

CREATE TABLE IF NOT EXISTS possession (
  id SERIAL,
  type INT NOT NULL,
  owner INT NOT NULL,
  price INT,
  PRIMARY KEY (id),
  FOREIGN KEY (owner) REFERENCES player(id),
  FOREIGN KEY (type) REFERENCES type(id)
);
```
- Create a new type of object or add new players,


```
INSERT INTO player (id, name, money)
VALUES
(DEFAULT, < name >, < money >);
```
- Change the name of a type of objects or the name of a player,


```
UPDATE player
SET
  name = < name >
```

```
WHERE
  id = < id >;
```

- Attribute an object of a given type to a player,


```
INSERT INTO possession (type, owner, price)
VALUES
(< typeId >, < ownerId >, NULL);
```
- Increase or decrease the amount of money a player has,


```
UPDATE player
SET
  money = money + < diff >
WHERE
  id = < id >;
```
- Retrieve the list of all the items that a player has,


```
SELECT
  type,
  owner
FROM
  possession
WHERE
  owner = < playerId >;
```
- Compute the current balance of a player,


```
SELECT
  money
FROM
  player
WHERE
  id = < id >;
```
- Allow a player to mark one of their item as buyable with a given price,


```
UPDATE possession
SET
  price = < price >
WHERE
  id = < id >
  AND owner = < playerId >;
```
- Allow a player to buy the cheapest item of a given type from the marketplace.


```
START TRANSACTION;

SELECT
  id,
  price,
  owner as curOwner
FROM
  possession
WHERE
  price IS NOT NULL
  AND type = < desired_type >
ORDER BY
  price ASC
LIMIT
  1;

UPDATE player
SET
  money = money + < price >
WHERE
```

```
id = curOwner >;
```

```
UPDATE player
SET
  money = money - < price >
WHERE
  id = < playerId >;

UPDATE possession
SET
  owner = < playerId >,
  price = NULL
WHERE
  id = < objectId >;

COMMIT;
```

Normalization & advanced exercises

- Create a department whose employees are located in different buildings using multivalued attributes.


```
CREATE TABLE IF NOT EXISTS department (
  dnumber INT NOT NULL,
  dname TEXT,
  d_head INT,
  d_building TEXT,
  PRIMARY KEY (dnumber)
);
```
- Retrieve information about a department based on the location,


```
SELECT
  *
FROM
  department
WHERE
  d_building LIKE '%BuildingB%';

SELECT
  *
FROM
  department
WHERE
  d_building = 'BuildingB';
```
- Normalize the department information to comply with 1NF.


```
CREATE TABLE IF NOT EXISTS department_normalized (
  id SERIAL,
  dnumber INT,
  dname TEXT,
  d_head INT,
  d_building TEXT,
  PRIMARY KEY (id)
);
```
- Retrieve information about a department based on the location using exact match.


```
SELECT
  *
FROM
  department_normalized
WHERE
  d_building = 'BuildingB';
```
- Create courses taught by the professors and attended by the students.

```
CREATE TABLE IF NOT EXISTS
```

```
teach (
  student TEXT,
  course TEXT,
  professor TEXT
);
```

- Define possible decompositions of the courses.

```
CREATE TABLE teach_1_1 AS
SELECT
  student,
  professor
FROM
  teach;
```

```
ALTER TABLE
  teach_1_1
ADD
```

```
PRIMARY KEY (student,
professor);
```

```
CREATE TABLE teach_1_2 AS
SELECT
  student,
  course
FROM
  teach;
```

```
ALTER TABLE
  teach_1_2
ADD
```

```
PRIMARY KEY (student,
course);
```

```
CREATE TABLE teach_2_1 AS
SELECT
  course,
  professor
FROM
  teach;
```

```
ALTER TABLE
  teach_2_1
ADD
```

```
PRIMARY KEY (professor);
```

```
CREATE TABLE teach_2_2 AS
SELECT
  course,
  student
FROM
  teach;
```

```
ALTER TABLE
  teach_2_2
ADD
```

```
PRIMARY KEY (course,
student);
```

```
CREATE TABLE teach_3_1 AS
SELECT
  course,
  professor
FROM
  teach;
```

```
ALTER TABLE
  teach_3_1
ADD
```

```
PRIMARY KEY (professor);
```

```
CREATE TABLE teach_3_2 AS
SELECT
  student,
  professor
FROM
  teach;
```

```
ALTER TABLE
```

```
teach_3_2
ADD
PRIMARY KEY (student,
professor);
```

- Reconstruct the courses taught by the professors and attended by the students.

```
SELECT
  course,
  professor,
  t1.student
FROM
  teach_1_1 AS t1,
  teach_1_2 AS t2
WHERE
  t1.student = t2.student;
```

```
SELECT
  t1.course,
  professor,
  student
FROM
  teach_2_1 AS t1,
  teach_2_2 AS t2
WHERE
  t1.course = t2.course;
```

```
SELECT
  course,
  t1.professor,
  student
FROM
  teach_3_1 AS t1,
  teach_3_2 AS t2
WHERE
  t1.professor =
t2.professor;
```

- Create information about specific employees by taking into account the properties of the employees in general.

```
CREATE TABLE IF NOT EXISTS employee (
  id SERIAL,
  name TEXT,
  salary INT,
  PRIMARY KEY (id)
);
```

```
CREATE TABLE IF NOT EXISTS professor (
  pid INT NOT NULL,
  field TEXT,
  PRIMARY KEY (pid)
) INHERITS (employee);
```

```
CREATE TABLE IF NOT EXISTS secretary (
  sid INT NOT NULL,
  building TEXT,
  PRIMARY KEY (sid)
) INHERITS (employee);
```

- Retrieve information about all the types of employees.

```
SELECT
  *
FROM
  employee;

SELECT
  *
FROM
  ONLY employee;

SELECT
```

```

*
FROM
  professor;

SELECT
  *
FROM
  secretary;

```

- Create courses having dependent types of information which would not exist otherwise.

```

CREATE TABLE IF NOT EXISTS
course (
  id SERIAL,
  name TEXT,
  PRIMARY KEY (id)
);

CREATE TABLE IF NOT EXISTS
session (
  course INT,
  num INT,
  name TEXT,
  PRIMARY KEY (course, num),
  FOREIGN KEY (course)
REFERENCES course(id)
);

```

TP4

Preliminary

Questions

- SeqScan: exploring the full table

```

SELECT
  *
FROM
  unicode;

```

- Index Scan: index to retrieve the position and fetch data

```

SELECT
  *
FROM
  unicode
WHERE
  codepoint = '0000';

```

- Index Only Scan: index to retrieve position

```

SELECT
  *
FROM
  unicode
WHERE
  codepoint < '0000';

```

- Bitmap Index Scan and Bitmap Heap Scan: bitmap using an index

```

SELECT
  *
FROM
  unicode
WHERE
  charname = '<control>';

```

- BitmapOr: two bitmaps created and then another bitmap is built with the OR condition

```

SELECT
  *
FROM
  unicode

```

```

WHERE
  numeric IS NOT NULL
or codepoint = '0000';

```

- BitmapAnd: two bitmaps created and then another bitmap is built with the AND condition

```

SELECT
  *
FROM
  unicode
WHERE
  numeric IS NOT NULL
AND charname < 'b';

```

- Filter: filter when no indexes are present

```

SELECT
  *
FROM
  unicode
WHERE
  comment IS NOT NULL;

```

- Nested Loop: join that cannot be efficiently processed

```

SELECT
  *
FROM
  unicode u1,
  unicode u2;

```

- Merge Join and Hash Join:

```

SELECT
  *
FROM
  unicode u1,
  unicode u2
WHERE
  u1.comment = u2.comment;

```

```

SELECT
  *
FROM
  unicode u1,
  unicode u2
WHERE
  u1.codepoint =
u2.lowercase;

```

TP5

Preliminary

Questions

- 1: Finds all the Shakespeare performances at Newcastle's Theatre Royal.

```

MATCH (theater:Venue
{name:'Theatre Royal'}),
(newcastle:City
{name:'Newcastle'}),
(bard:Author
{lastname:'Shakespeare'}),
(newcastle)<-
[:STREET|CITY*1..2]-
(theater)
<-[:VENUE]-()-
[:PERFORMANCE_OF]->()
-[:PRODUCTION_OF]-
>(play)<-[:WROTE_PLAY]-
(bard)
RETURN DISTINCT play.title
AS play

```

- 2: Finds all the Shakespeare performances at Newcastle's Theatre Royal after 1608.

```

MATCH (theater:Venue
{name:'Theatre Royal'}),
(newcastle:City
{name:'Newcastle'}),
(bard:Author
{lastname:'Shakespeare'}),
(newcastle)<-
[:STREET|CITY*1..2]-
(theater)
<-[:VENUE]-()-
[:PERFORMANCE_OF]->()
-[:PRODUCTION_OF]-
>(play)<-[:WROTE_PLAY]-
(bard)
WHERE w.year > 1608
RETURN DISTINCT play.title
AS play

```

- 3: How many Shakespeare performances were at Newcastle's Theatre Royal?

```

MATCH (theater:Venue
{name:'Theatre Royal'}),
(newcastle:City
{name:'Newcastle'}),
(bard:Author
{lastname:'Shakespeare'}),
(newcastle)<-
[:STREET|CITY*1..2]-
(theater)
<-[:VENUE]-()-
[:PERFORMANCE_OF]->()
-[:PRODUCTION_OF]-
>(play)<-[:WROTE_PLAY]-
(bard)
RETURN count(play)

```

- 4: Rank plays by number of performances.

```

MATCH (theater:Venue
{name:'Theatre Royal'}),
(newcastle:City
{name:'Newcastle'}),
(bard:Author
{lastname:'Shakespeare'}),
(newcastle)<-
[:STREET|CITY*1..2]-
(theater)
<-[:VENUE]-()-
[p:PERFORMANCE_OF]->()
-[:PRODUCTION_OF]-
>(play)<-[:WROTE_PLAY]-
(bard)
RETURN play.title AS play,
count(p) AS
performance_count
ORDER BY performance_count
DESC

```

- 5: Find the plays written by Shakespeare, and order them based on the year in which they were written. (HINT: Use WITH and collect())

```

MATCH (bard:Author
{lastname:'Shakespeare'})-
[w:WROTE_PLAY]->(play)
WITH play
ORDER BY w.year DESC
RETURN collect(play.title)
AS plays

```